



0 E T 2 0 7

How Swiggy uses Amazon ElastiCache to power their ML feature store

Smita Srivastava

Solutions Architect
AWS India

Shri Nath Agarwal

Software Developer IV
Swiggy



Agenda

- Quick overview of AWS databases
- Amazon ElastiCache use cases
- Machine Learning(ML) process flow and need for a feature store
- Benefits of Amazon ElastiCache as an ML feature store
- About Swiggy and its ML landscape
- Swiggy's requirements from a feature store
- Why Amazon ElastiCache was a great fit
- Resources

Most complete set of relational & purpose-built databases



Amazon
RDS



Amazon
Aurora

KEY-VALUE



Amazon DynamoDB

DOCUMENT



Amazon DocumentDB

CACHING



Amazon ElastiCache

GRAPH



Amazon Neptune

TIME-SERIES



Amazon Timestream

MEMORY



Amazon MemoryDB

WIDE COLUMN



Amazon Keyspaces

LEDGER



Amazon QLDB

Amazon ElastiCache use cases




Real-time application data caching

Store frequently used data in-memory for microsecond response times and high throughput



Session store

Store ephemeral session data (authentication information, cookies) for internet-scale applications



Cost-optimized real-time analytics performance

Reduce cost and access real-time analytics performance



Geospatial

Real-time geospatial data at scale and speed, for ride-hailing, online dating, social media, and gaming



Leaderboards

Simplify application development with built-in data structures including re-ranking of sorted set in real time



Media streaming

Store live streaming data including metadata for user profiles, authentication



Chat and messaging

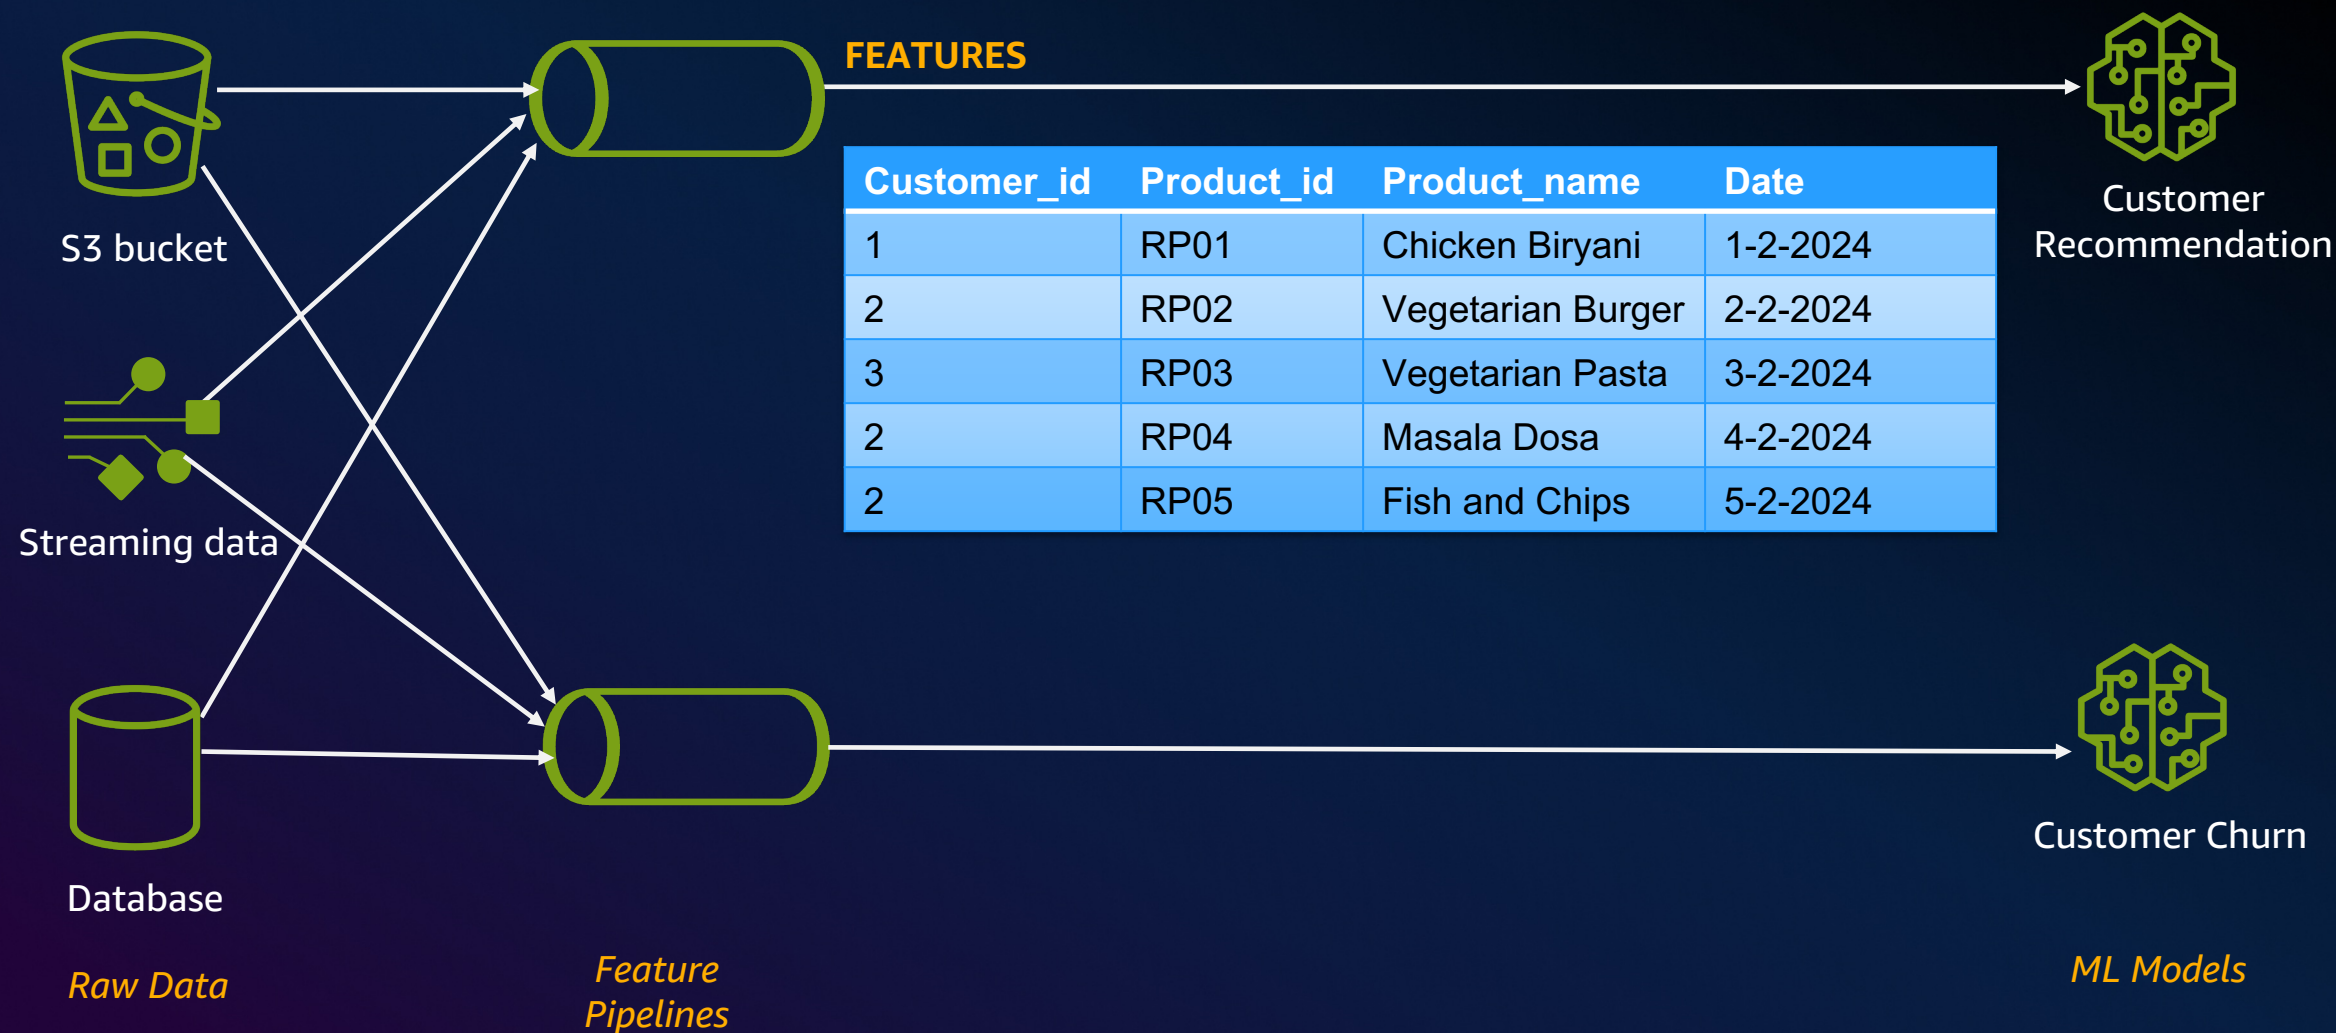
Use PUB/SUB standard to support high performance chat rooms and real-time comments



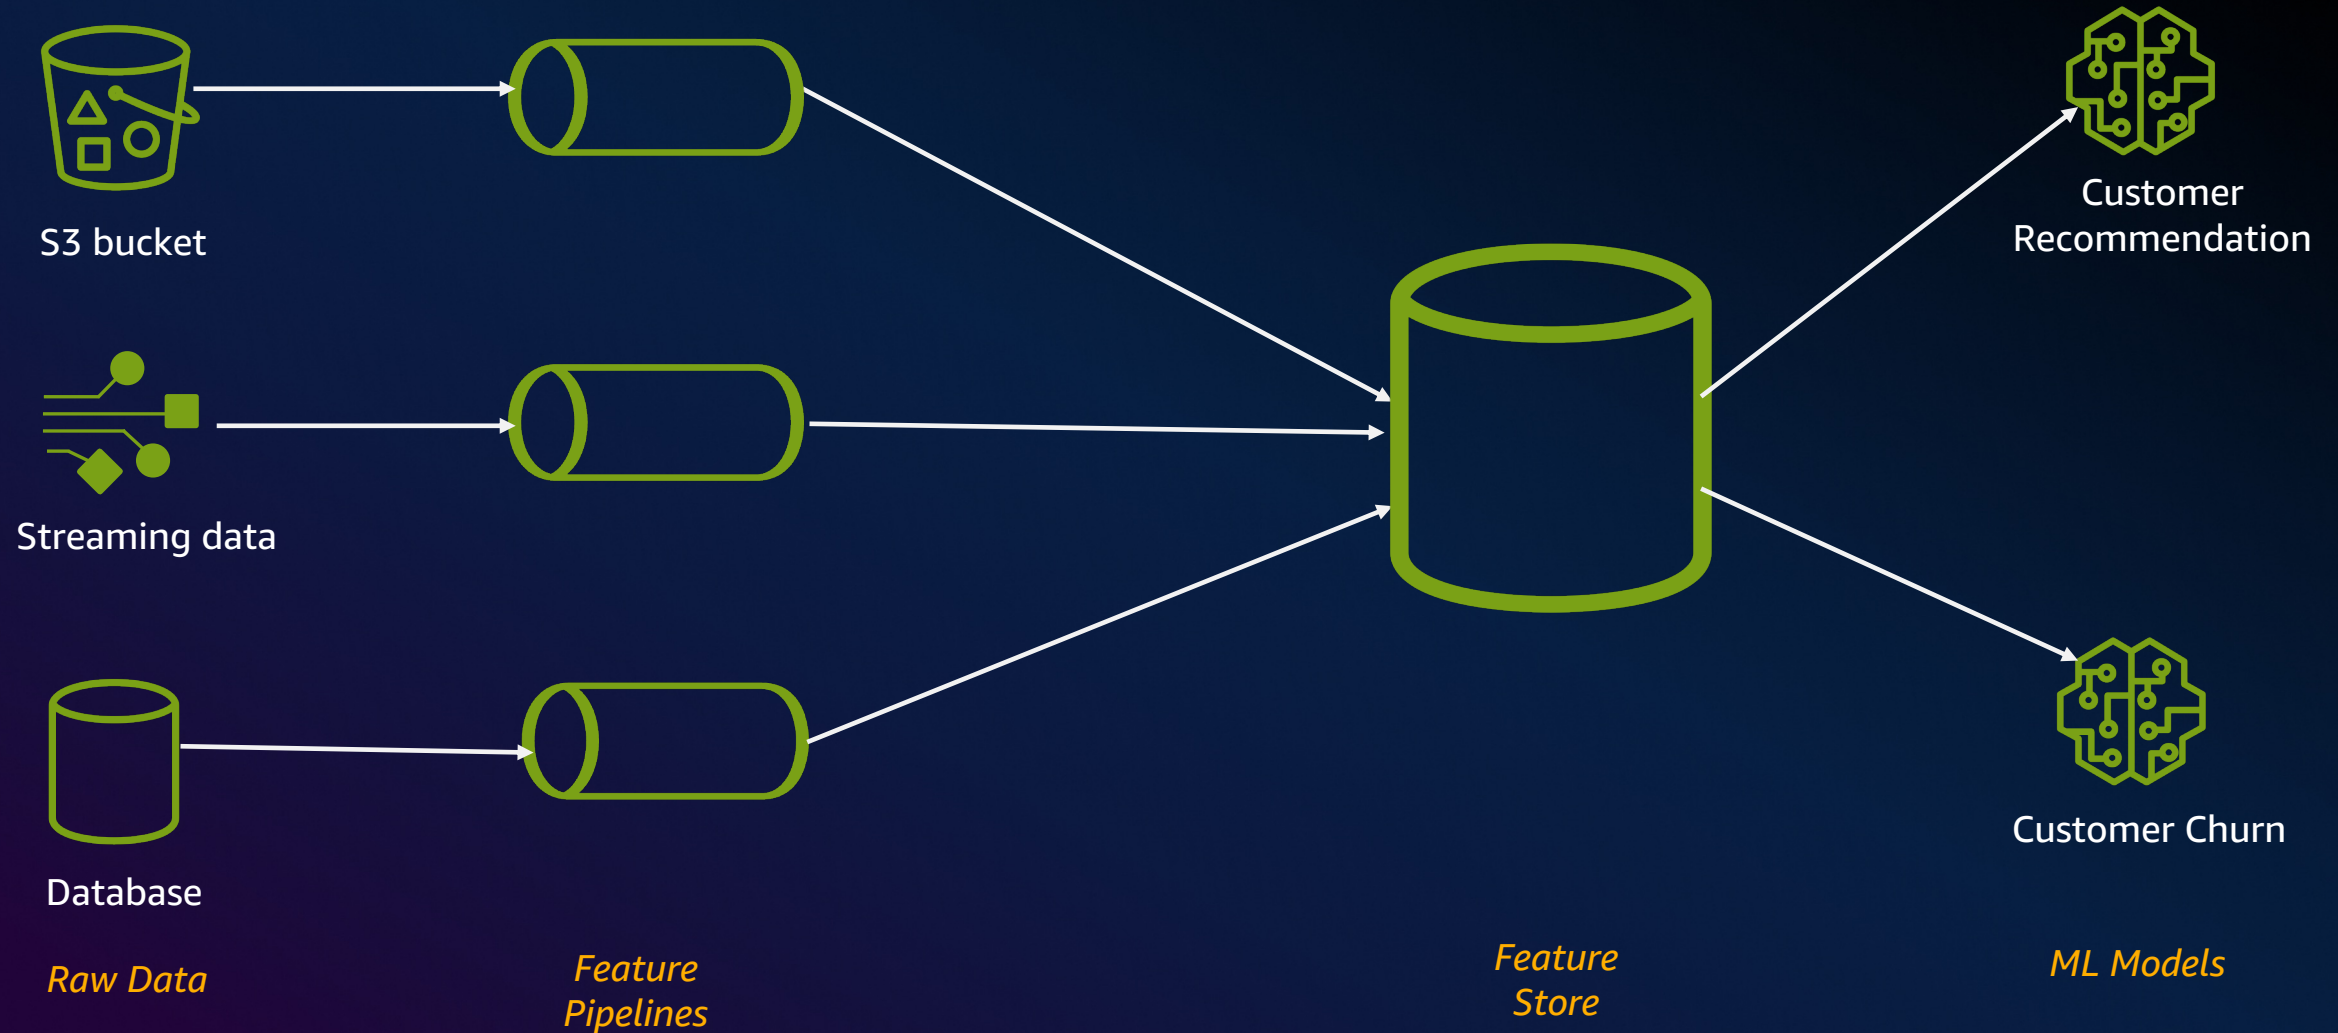
Machine learning

Feature store for machine learning workloads

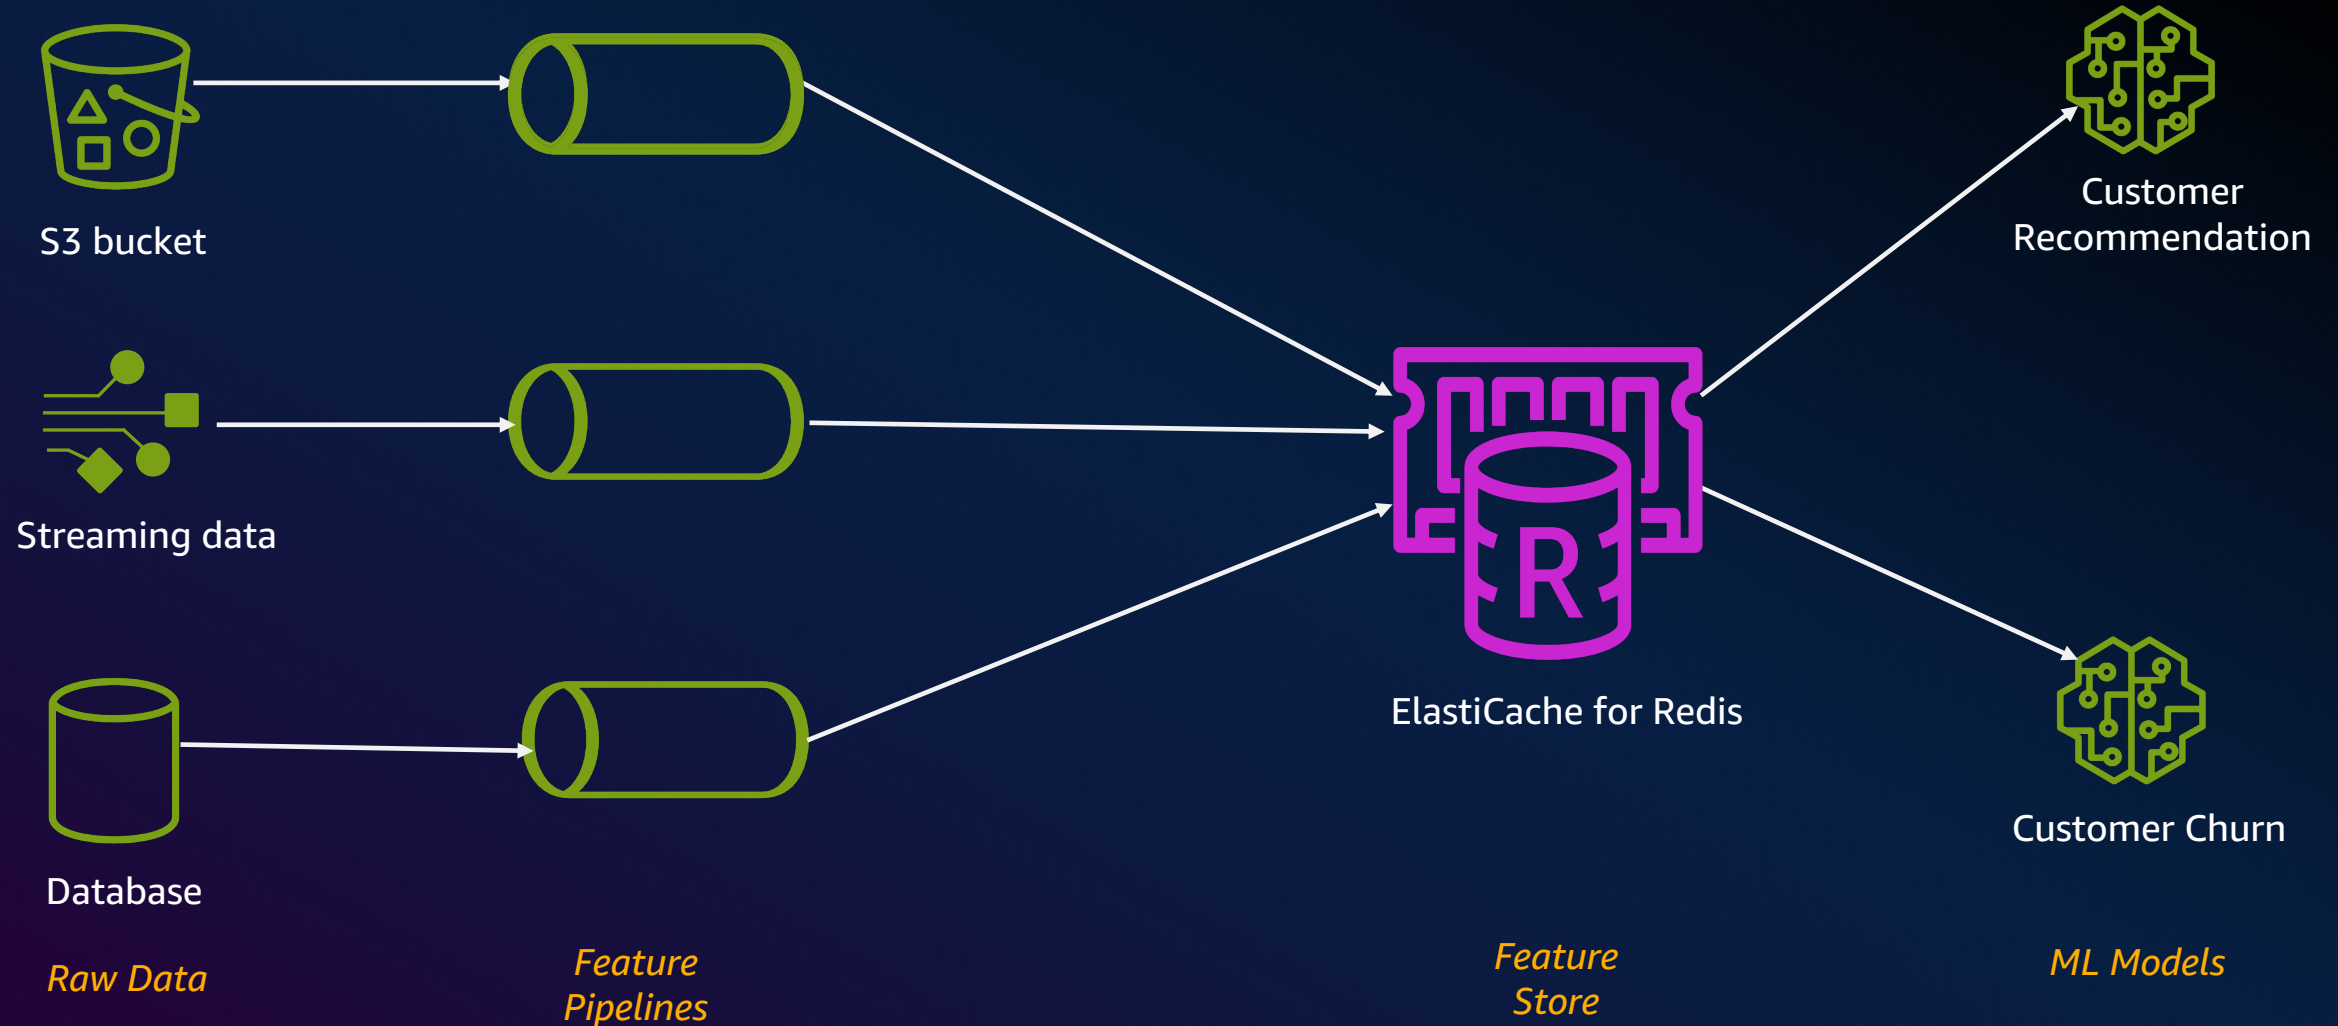
Typical Machine Learning(ML) workflow



Machine Learning(ML) workflow with a Feature Store



ML process with Amazon ElastiCache(for Redis) as an Online Feature Store



Benefits of Amazon ElastiCache(EC) as a Feature Store

- High Performance

In-memory caching, scales to millions of operations per second, sub millisecond read/write response times

- Fully Managed

Hardware provisioning, software patching, configuration, monitoring, failure recovery, backups

- Highly Available and Reliable

EC clusters deployed across multiple AZs with automatic failover

- Cost Effective

Pay as you go pricing

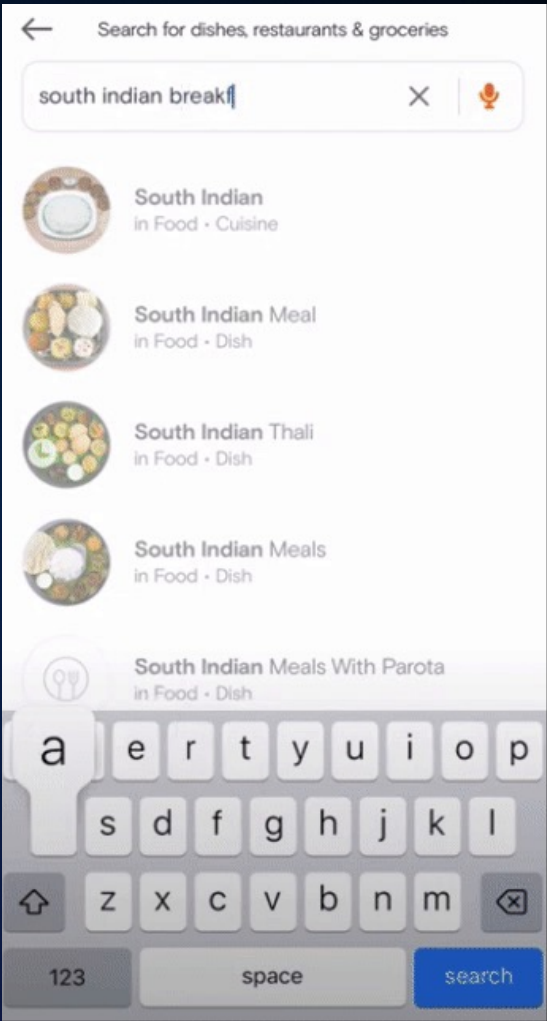
- Scalable architecture

Designed to support online cluster resizing

- Security Controls

IAM support, encryption using customer managed keys stored in KMS, secure internode communications

Swiggy – 3 Way Market Place



100+ ML Models



1 million prediction at peak

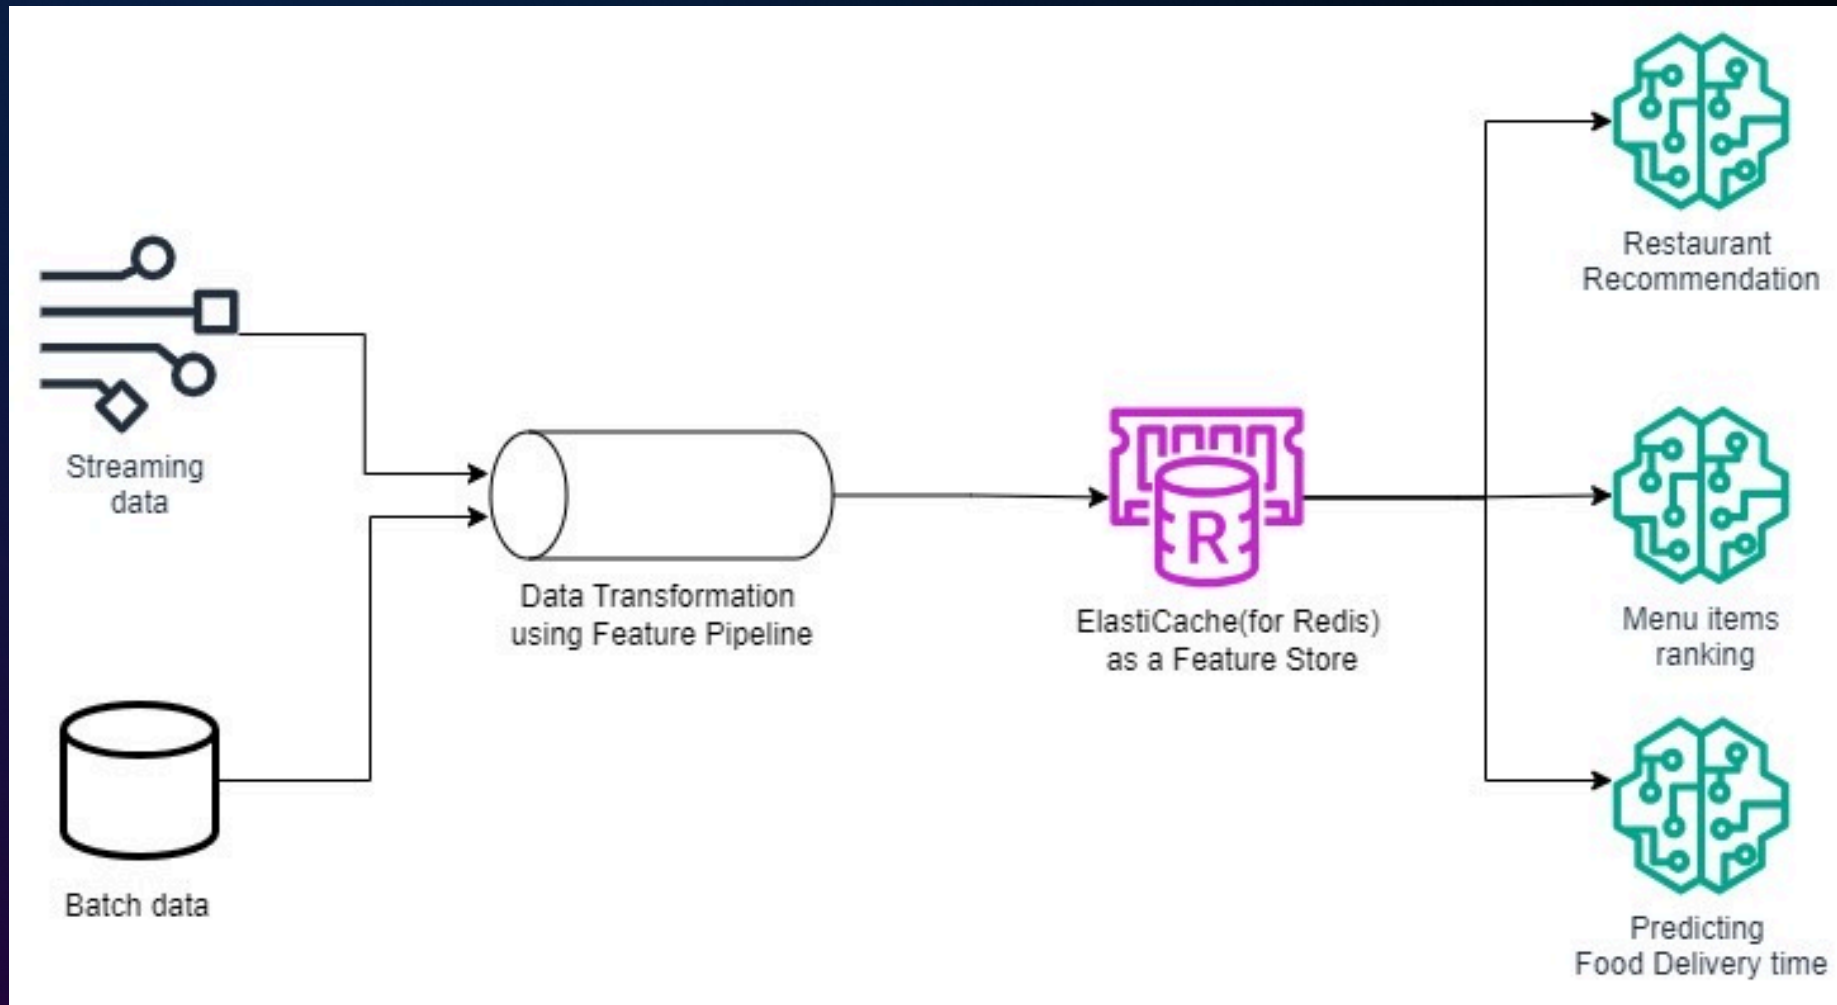


1000+ ML Features

Swiggy's requirements from an ML Feature Store

- **Persistent scalable storage** – handling billions of Records to be stored
- **High Read Throughput** – supporting millions of Feature look up per second
- **Autoscaling** – on-demand scaling
- **Efficient Batch Writes** – enabling full data refresh
- **Low latency** – sub milliseconds to fetch the features
- **Cost Effective**

Amazon ElastiCache for Redis as a feature store for Swiggy



Benefits of Amazon ElastiCache as a Feature Store for Swiggy

- millions/second read requests served from ElastiCache
- <10 microseconds response time
- Reduced management effort
- Automatic backup used for analytics workload
- Read replicas to achieve higher availability by spreading out read workloads
- No manual efforts in autoscaling
- Auto-failover - no midnight pager alerts when a node goes down.
- Cost effective

Amazon ElastiCache Resources



Browse Amazon
ElastiCache
Resources



ElastiCache as a
Feature store

skillbuilder.aws 

Build beyond

Redeem your free 7-day
trial of AWS Skill Builder

Thank you!

Smita Srivastava

Solutions Architect
AWS India

Shri Nath Agarwal

Software Developer IV
Swiggy



Please complete the session
survey in the mobile app