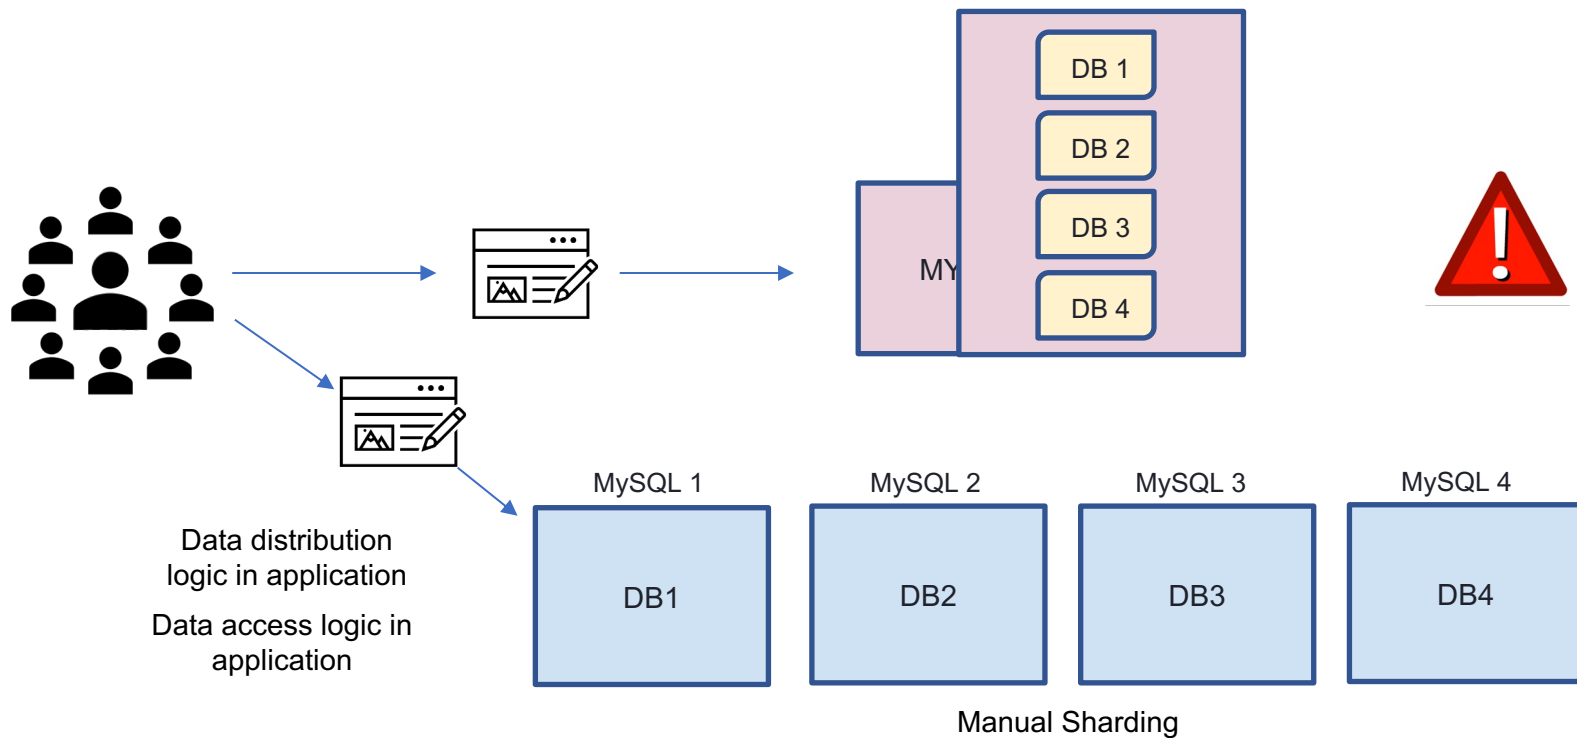




SUMMIT
INDIA

From data to discovery: Exploring the intersection of distributed database and AI

Data Growth: A challenge



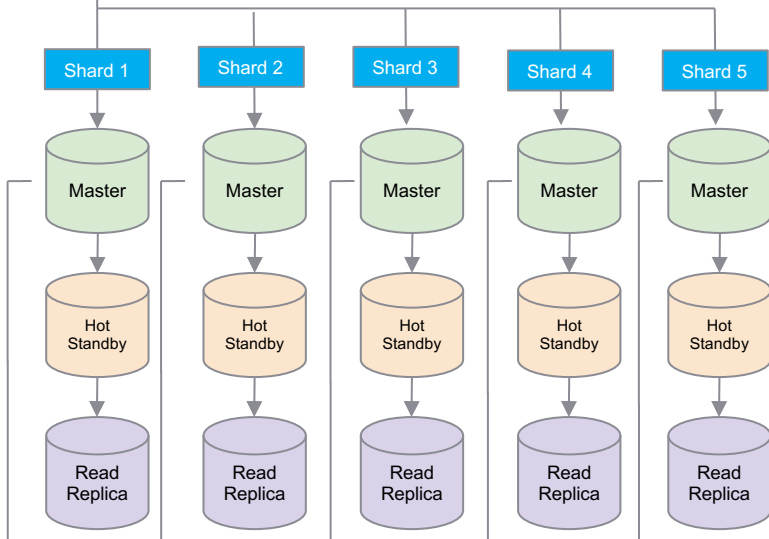
Intent to grow creates Complexity

Application **Region 1 - Active**



Reads

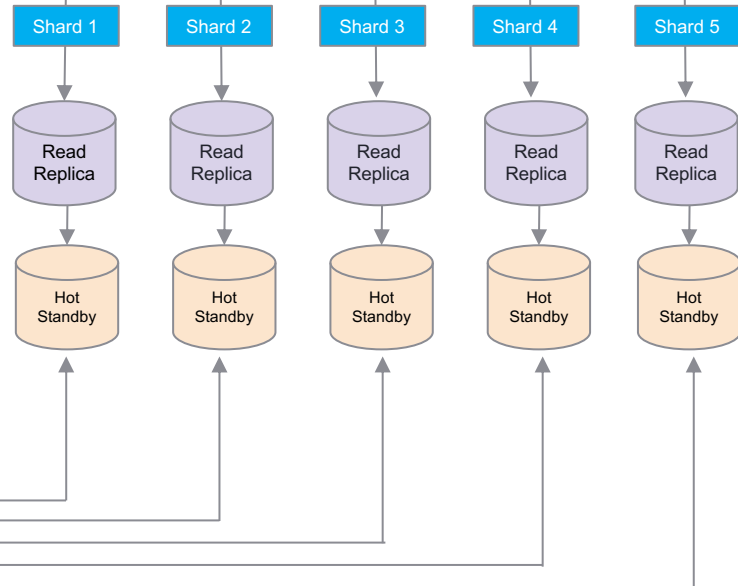
Writes



Application **Region 2 - Passive**



Reads



25 DB Nodes

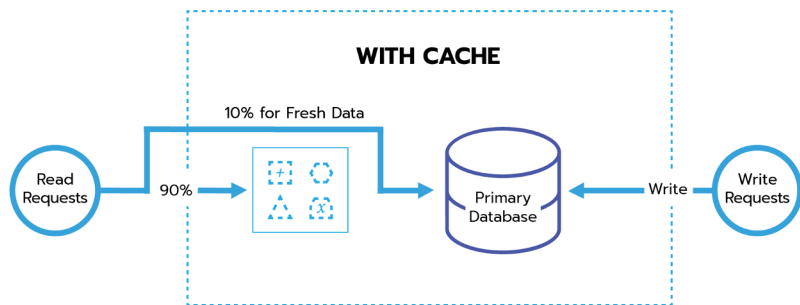
20 Replication Channels

Operational complexity

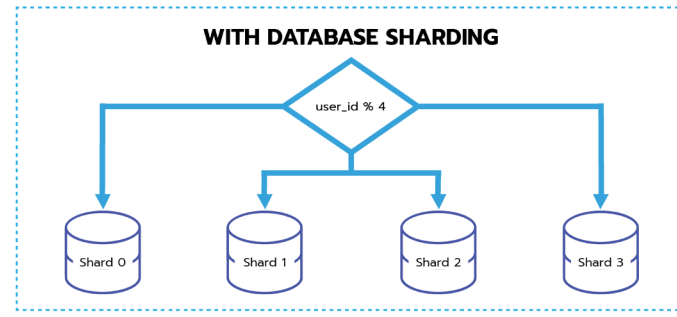
Failover complexity

Resharding complexity

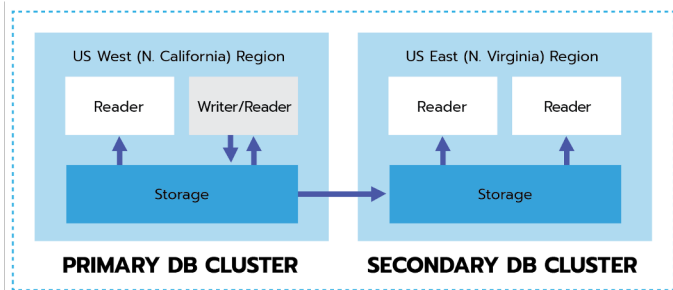
Ways to Scale



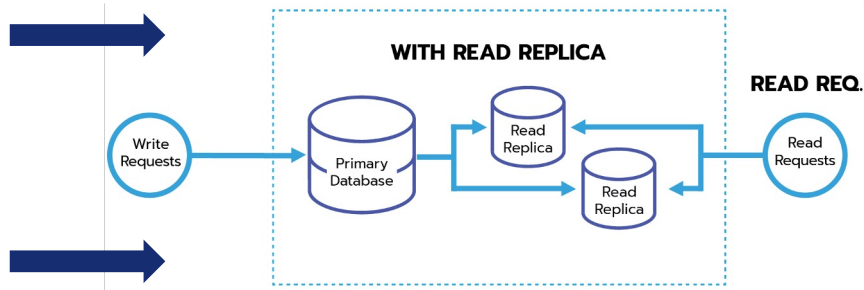
Caching Layer



Sharding




Primary Secondary Replication



Separated Read-Write

Single Node Database Limitations



Scalability

Unstable performance when scaling write-intensive applications.




High Availability

Setting up HA requires careful planning & configuration. The replication can result in lag, potential data inconsistencies



Real-Time Analytics

Analytical queries can impact transactional processing




Handle Modern applications

Adapting to cloud native architecture poses challenges to traditional, monolithic systems like MySQL

Introducing TiDB

The most advanced,
open source, distributed
SQL database for
modern applications.

Scalable. Versatile. Titanium (Ti) Reliable.



34K+ GitHub Stars



3K+ Adopters

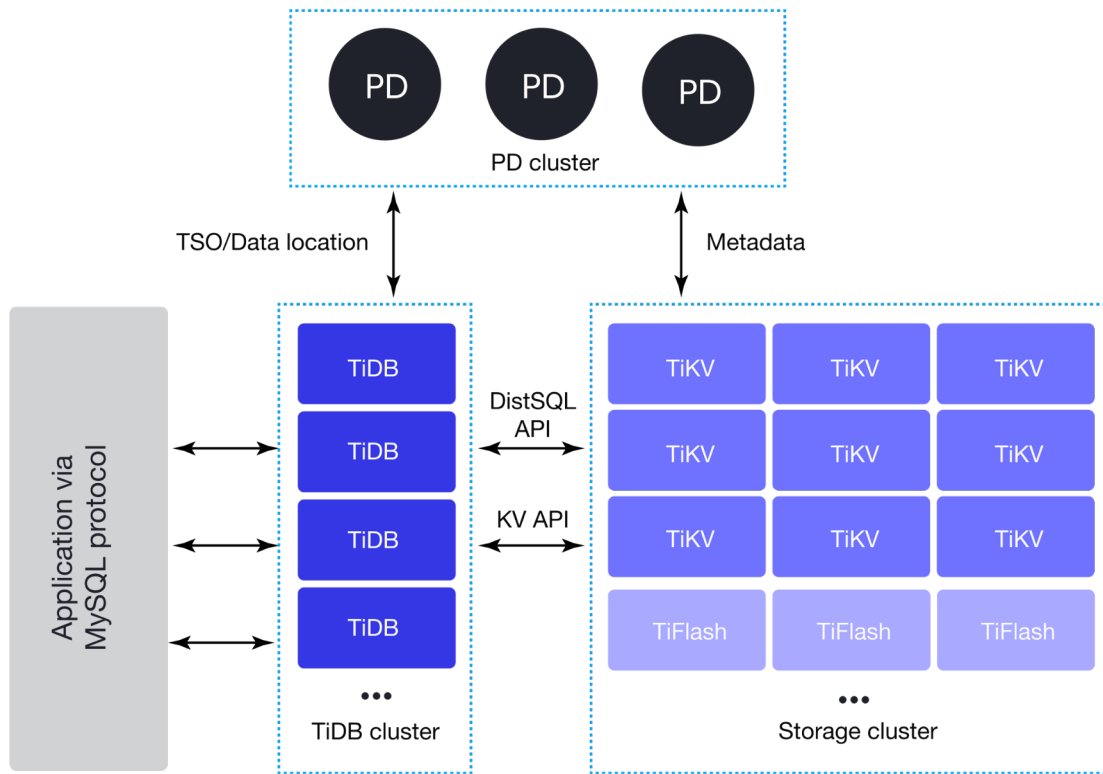


800+ Contributors




5K+ Slack Users

What Makes TiDB So Advanced?



The Advantages of TiDB



Horizontal Scaling

Grants total transparency into data workloads without manual sharding.




High Availability

Guarantees auto-failover and self-healing for continuous access to data.



Mixed Workloads

Streamlined tech stack makes it easier to produce real-time analytics.



MySQL Compatibility

Enjoy the most MySQL compatible distributed SQL database on the planet.



Multi-Cloud

Deploy database clusters anywhere in the world.



Open Source


Unlock business innovation with a database that's 100% open source.



Robust Security


Protect data with enterprise-grade encryption both in-flight and at-rest.

Top Use Cases for TiDB




MySQL Alternative

Migrate to a more affordable and elastic MySQL alternative that supports real-time analytics right out of the box.




Real-Time Analytics

Enable your business to process and query new data as it's created to guide decision making, enhance resource utilization, and improve customer experiences.



Application Modernization

Boost developer productivity with a modern, distributed SQL database that offers true elastic scale and relentless reliability combined with mixed workload processing.



Tech Stack Unification

Reduce costs and system complexity with a unified data stack that can replace traditional relational databases, NoSQL databases, and lightweight data warehouses.

Target Verticals: Suitability and Why

Fintechs especially Crypto / P2P	Gaming especially Social Gaming	Logistic	SaaS	E-Commerce
Suitability: ★★★★★	Suitability: ★★★★★	Suitability: ★★★★★	Suitability: ★★★★★	Suitability: ★★★★★
<ul style="list-style-type: none"> • TiDB is battle tested in money related scenarios due to high resilience and consistency • Very simple architecture for scalability and real-time insights • Usually they don't pay much attention building tech team as other verticals 	<ul style="list-style-type: none"> • Fast pace and release new game titles very often (mobile gaming) • Cares a lot about real-time insights • Need to be prepared for user growth • Used a lot in billing 	<ul style="list-style-type: none"> • Cares a lot about real-time insights • Top players have large volume of data • Engineering teams' skill level is slightly weaker than DNB companies with same scale 	<ul style="list-style-type: none"> • Scalability and multi-tenancy • Need either gather user data from multiple sources into one or need to serve all users in fewer number of database instances to decrease maintenance cost • Operational dashboard / analytics is a must for most of them 	<ul style="list-style-type: none"> • Scalability is one of the most important things for them • Better performance is better customer experience for them • Various new scenarios everyday cause trouble in maintaining and GTM
China Construction Bank PINGAN Group Paypay Square-Cash App	NetEase Games GAEA Colopl Kunlun	ZTO Express Ninja Van Dehlivery	Patsnap Streak Databricks	Flipkart Shopee

Scaling TiDB To 1 Million QPS

A Flipkart journey



Sachin Japate · Follow

Published in Flipkart Tech Blog · 9 min read · Jul 10



304



1



Mammoth Performance
in Throughput & Latency

1,070,000

ops per second (consistent reads)
Latency: **4.82ms** P95 | **7.43ms** P99


123,000

ops per second (writes)
Latency: **6.1ms** P95 | **13.1ms** P99

Real Numbers from Flipkart!

Vectorized Query Execution



 Sunny Bains · 1st
swe@PingCAP - the company behind TiDB.
1d · 🌐

This is just too cool 😍. Exciting development, stay tuned!


```
tidb> CREATE TABLE my_vec_table(val VECTOR);
Query OK, 0 rows affected (0.16 sec)

tidb> INSERT INTO my_vec_table VALUES
-> ('[8.7, 5.7, 7.7, 9.8, 1.5]'),
-> ('[3.6, 9.7, 2.4, 6.6, 4.9]'),
-> ('[4.7, 4.9, 2.6, 5.2, 7.4]'),
-> ('[7.7, 6.7, 8.3, 7.8, 5.7]'),
-> ('[1.4, 4.5, 8.5, 7.7, 6.2]');
Query OK, 5 rows affected (0.05 sec)
Records: 5 Duplicates: 0 Warnings: 0

tidb> -- Select closest vectors
Query OK, 5 rows affected (0.03 sec)
Records: 5 Duplicates: 0 Warnings: 0

tidb> SELECT val,
-> VEC_Cosine_Distance(val, '[1,2,3,4,5]') AS d
-> FROM my_vec_table
-> ORDER BY d;
+-----+-----+
| val                | d                |
+-----+-----+
| [1.4,4.5,8.5,7.7,6.2] | 0.04972827044697803 |
| [4.7,4.9,2.6,5.2,7.4] | 0.06925255631855243 |
| [7.7,6.7,8.3,7.8,5.7] | 0.12677426832319694 |
| [3.6,9.7,2.4,6.6,4.9] | 0.18577333207371582 |
| [8.7,5.7,7.7,9.8,1.5] | 0.2564100235012563 |
+-----+-----+
5 rows in set (0.03 sec)

tidb> █
```

 Ed Huang @dxhuang · 18h
Looks pretty cool :)

8 3 comments · 1 repost

Vector databases make it easier for machine learning models to remember previous inputs, allowing machine learning to be used to power search, recommendations, and text generation use-cases.

Some Vector Use Cases



Recommendation Engines

Used by E-Commerce, content streaming platforms, personalized marketing. Vector offers similarity calculations



NLP Processing


Sentiment Analysis, text classification, documents similarity. Facilitates efficient text search, clustering



Fraud Detection

Used in cybersecurity, financial services, fraud prevention. This is possible with pattern recognition & similarity search


Embedding and Vector




Distributed TiDB and Vector together!

MySQL+Vector DB(multiple round trip) -> TiDB(single round trip)

Complex Architecture



Simple Architecture



Trusted by Global Innovation Leaders

Global Top Lighthouses



Top Fintech in the World



E-commerce and DNB Lighthouses






TiDB Cloud

is hosted on 


TiDB Serverless *Hosted on AWS*

The Revolutionary Serverless Database

<https://tidbcloud.com>



*Every developer has a **FREE & scalable** database*



Autopilot Scaling



Zero Downtime



Cost Efficiency



Easy to Use

Thank You!