




# Redis data integration: Cost-effective performance at scale

**Shekhar Suman,**  
Redis Multi-modeling Data Architect, Redis

**Abhishek Srivastava,**  
Redis Data Integration Guru, Redis



# The race for digital transformation and being relevant



## Declining top-line growth

Financial firms are experiencing declining top-line growth due to low interest rates and new regulations



## Changing expectations

Consumers and clients are demanding responsive, always-on, omni-channel experiences



## Industry disruptors

The pandemic led to between a 21% and 26% increase in the relative rate of daily downloads of Fintech apps.



## Availability drives revenue


E-commerce companies seeing revenue loss for each second their applications are down

# Problems with legacy systems



## Slow

Legacy systems are slow and unable to keep up with real-time needs.




## Fragmented

Separate point solutions have resulted in fragmented data silos.



## Not resilient

Business disruption and growing digital demands are resulting in systems failures.



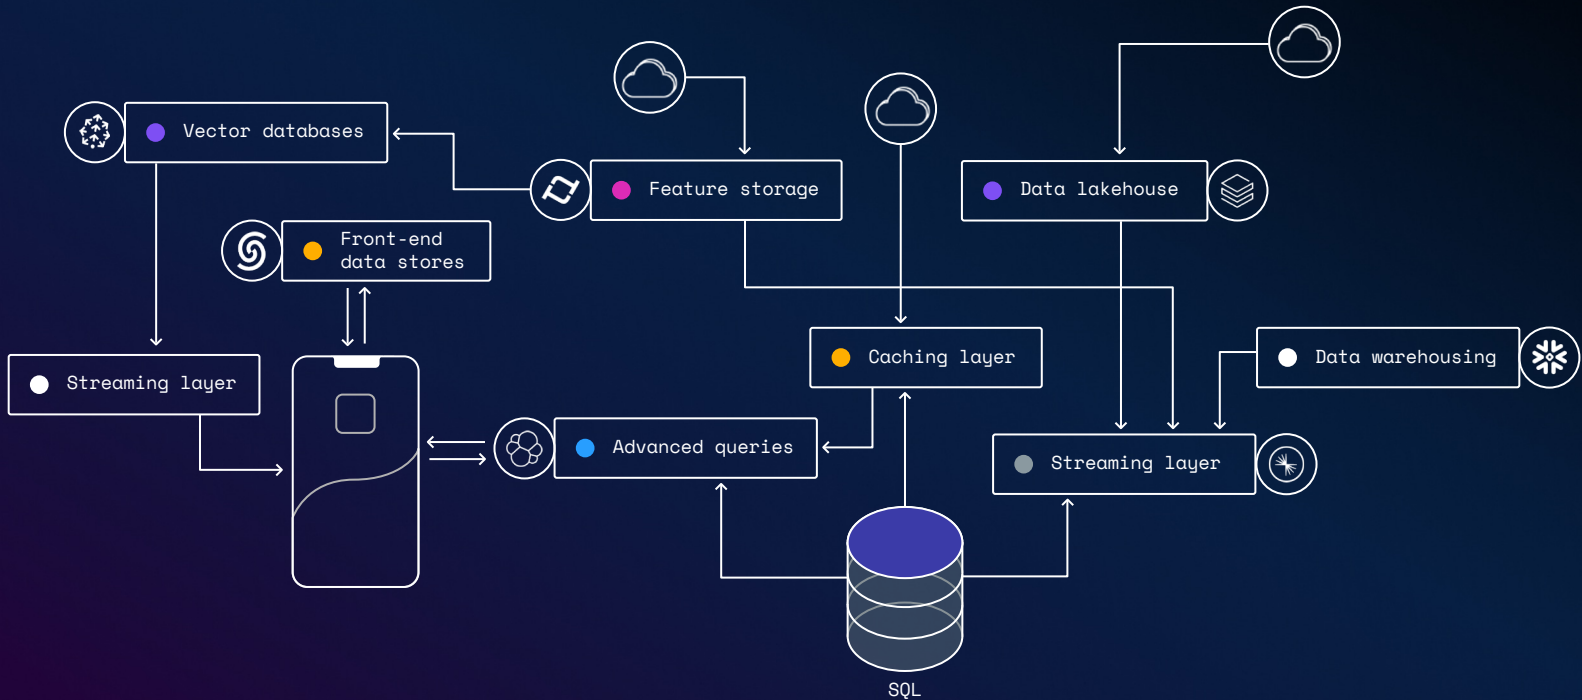
## Inefficient

Legacy systems hinder cost efficiency due to expensive technologies and talent shortage.



“Organisations need an efficient, cost-effective and proven technology thereby boosting their team’s productivity and addressing most of their technical problems”

# Applications are becoming more complex



# The Problem Statement

- **Have fleet of different traditional databases. Hence performance and scale are limited**
- **Business won't scale due to limitations of underlying databases**
- **Application is critical - application is real time - Need faster time to production**
- **Improve overall user experience with consistent performance**
- **Standards are high and market is tough - want to stay ahead in competition**
- **Need a technology that provides availability SLA guarantees & linearly scalable**
- **TCO and maintainability are important**





# How do we solve this problem?

“We solve this problem by bringing in a technology that not only can serve as a high-throughput database but can also simplify application architectures and enhance operational efficiency.

This technology should fit harmoniously with traditional databases to facilitate a seamless transition“



# Let's understand by building a true real-time app



**A sample securities trading and account portfolio application**



# Building Blocks

## Redis Data Integration\*

Data Pipeline between Legacy systems and Redis



### Ingest

Works as tool to ingest data from various sources to Redis



### Write Behind

Works as a tool to write data from Redis down to various databases



### Configuration Driven

No Code required, works completely in configurations provided in yaml files



### Multi database Support

Most of the databases are supported out of the box and new features being added continuously



## Core Components

Part of Redis enterprise



### Streams

High throughput streaming system which supports persistence and consumer groups



### Search

Complex Queries on JSON data along with filters aggregations and Full text Search



### Time Series

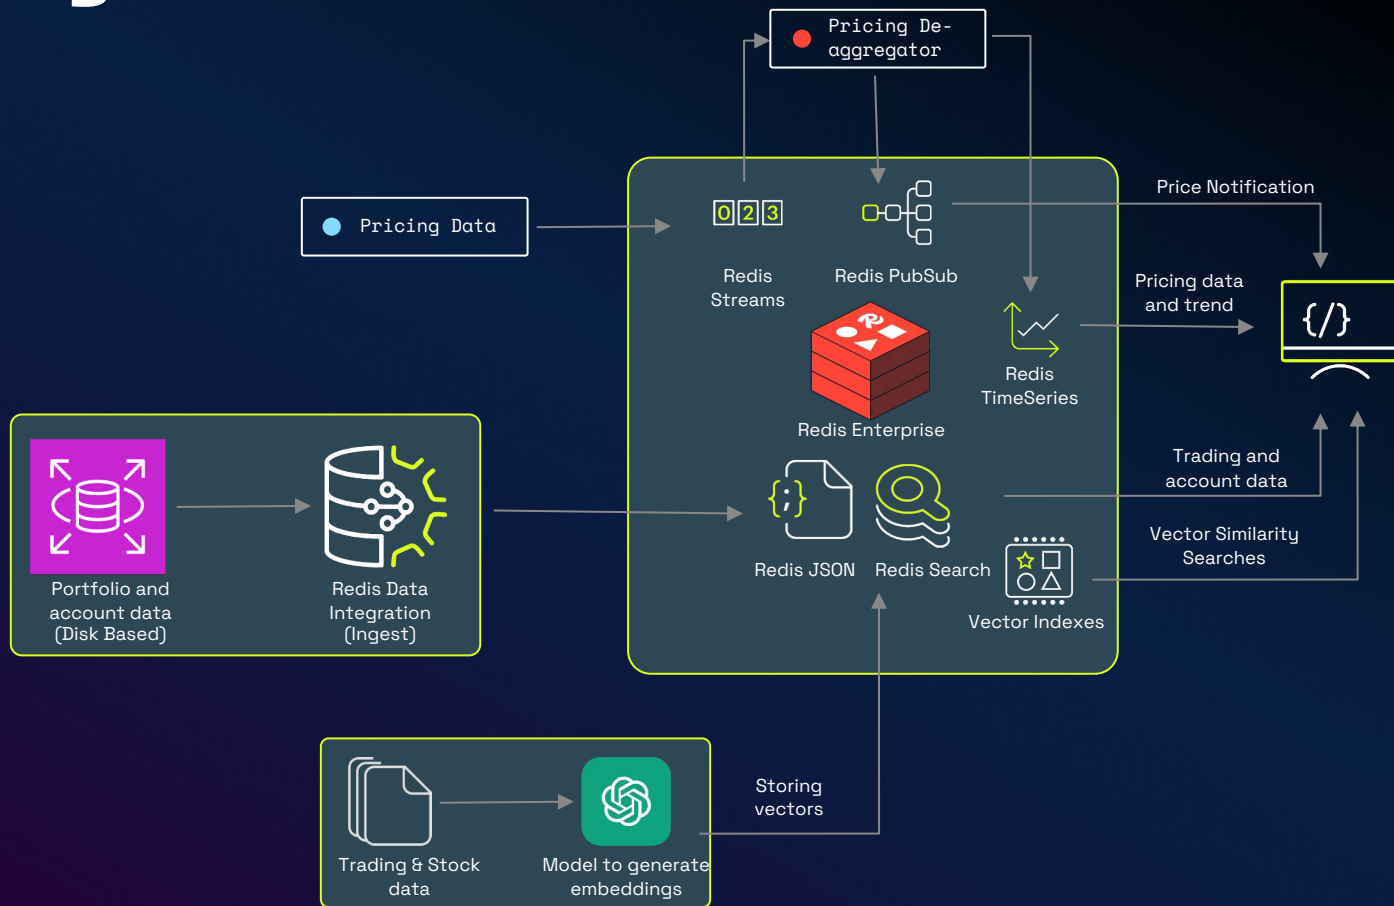
Fast timeseries storage with supports downsampling out of the box



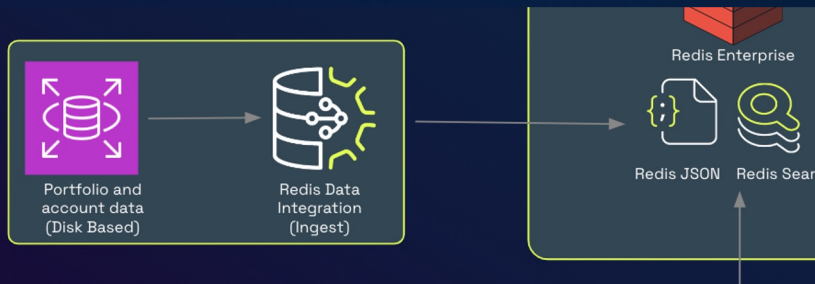
### JSON

High performance NoSQL Document storage which can be scaled to very high throughputs

# Redesigned Architecture



# Ingest Account and security data



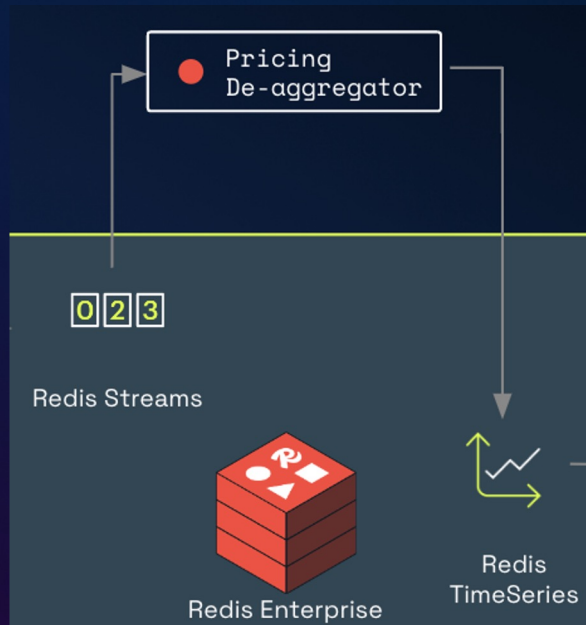
- MySQL (RDS) contains account and trading data
- Use RDI to ingest the data from MySQL to Redis Enterprise
- Data in Redis is denormalized as a JSON documents which will be indexed using Redis Search capability
- The user can query this data for analytical and reporting purposes

# Generate pricing data



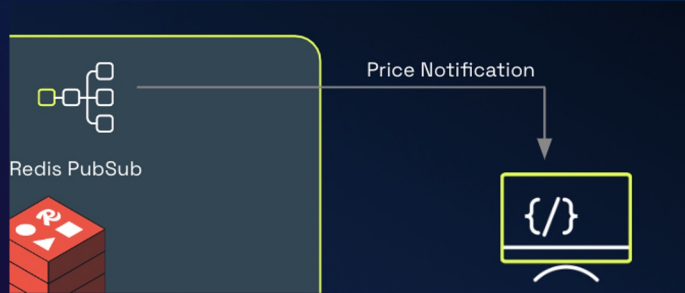
- High speed stock ticker data is ingested into Redis Streams topic in Redis Enterprise
- The consumer will process these ticker data asynchronously

# Consume & deaggregate pricing data



- Pricing De-aggregator acts as a consumer to the Redis Streams topic and processes these pricing data
- The pricing data is then converted into Timeseries data and can be down-sampled

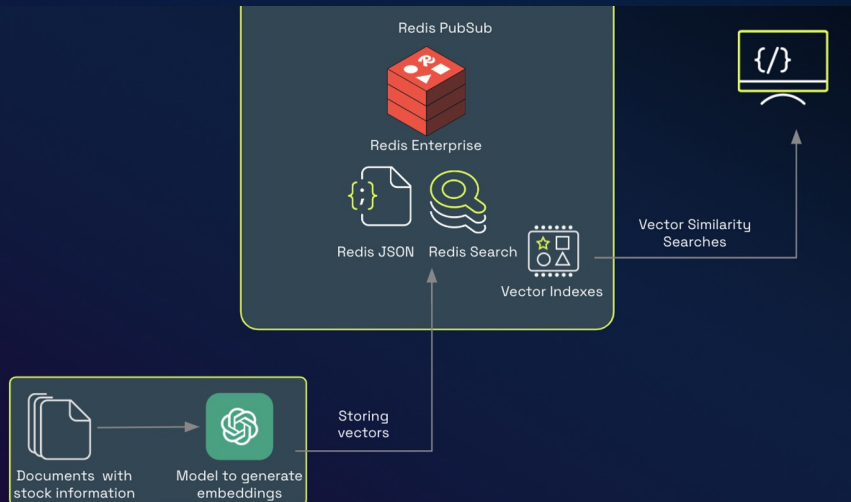
# Send pricing notifications



- The De-aggregated data is also sent to redis pubsub which can trigger a price notification to the user

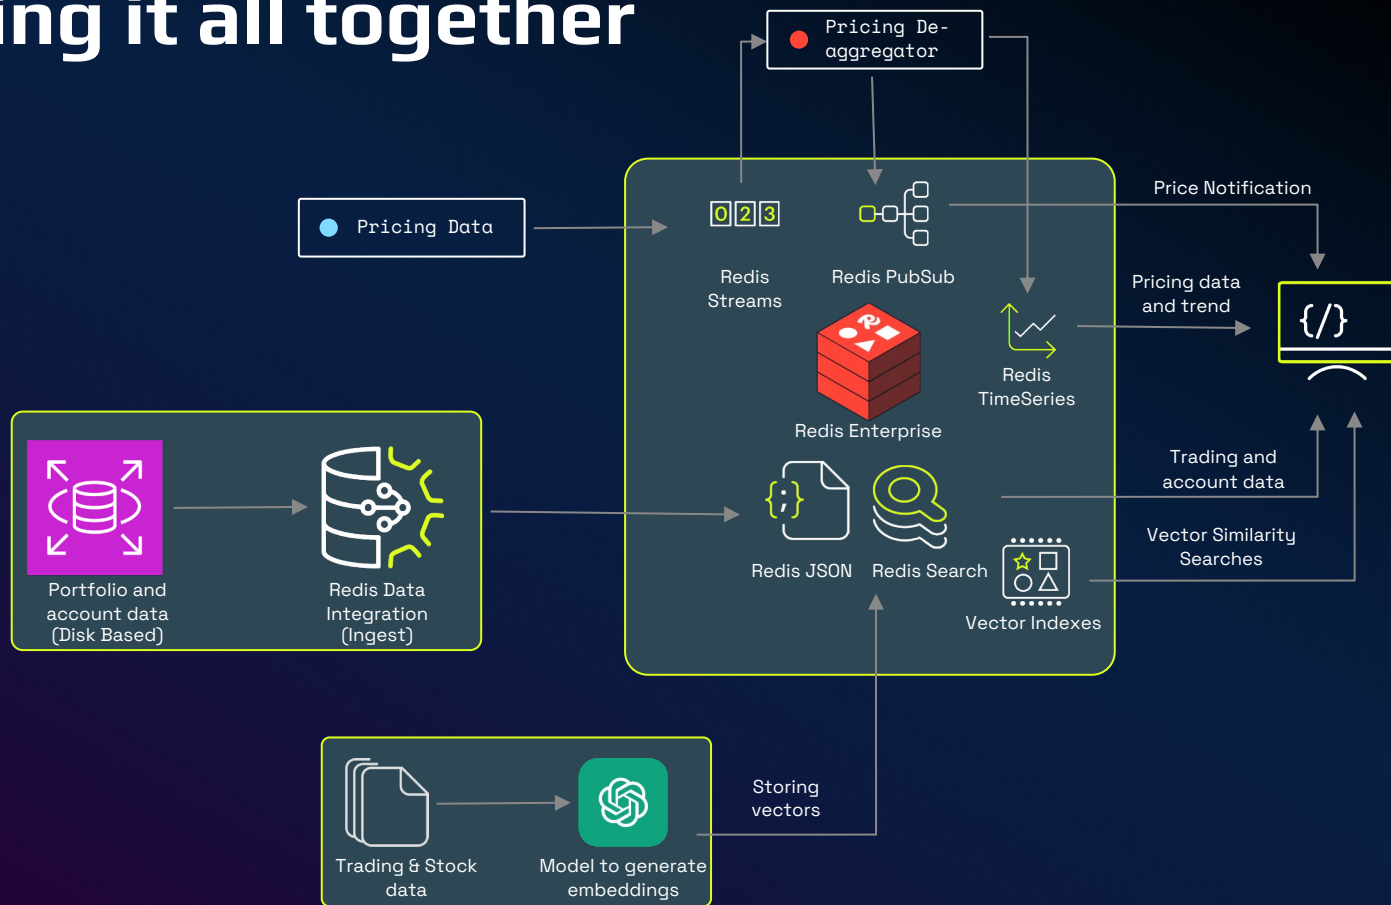


# Perform semantic searches on related docs



- Documents are passed through a model to generate embeddings
- These embeddings are then stored in Redis
- Appropriate indexes are created on the vectors along with other data
- Vector Similarity searches are performed to do semantic and hybrid searches on the indexes

# Putting it all together



# Go beyond caching



Enterprise  
caching



Primary  
database



Session  
management



Messaging



Fraud  
detection



Real-time  
inventory



Leaderboards



Vector  
database



Online  
feature store

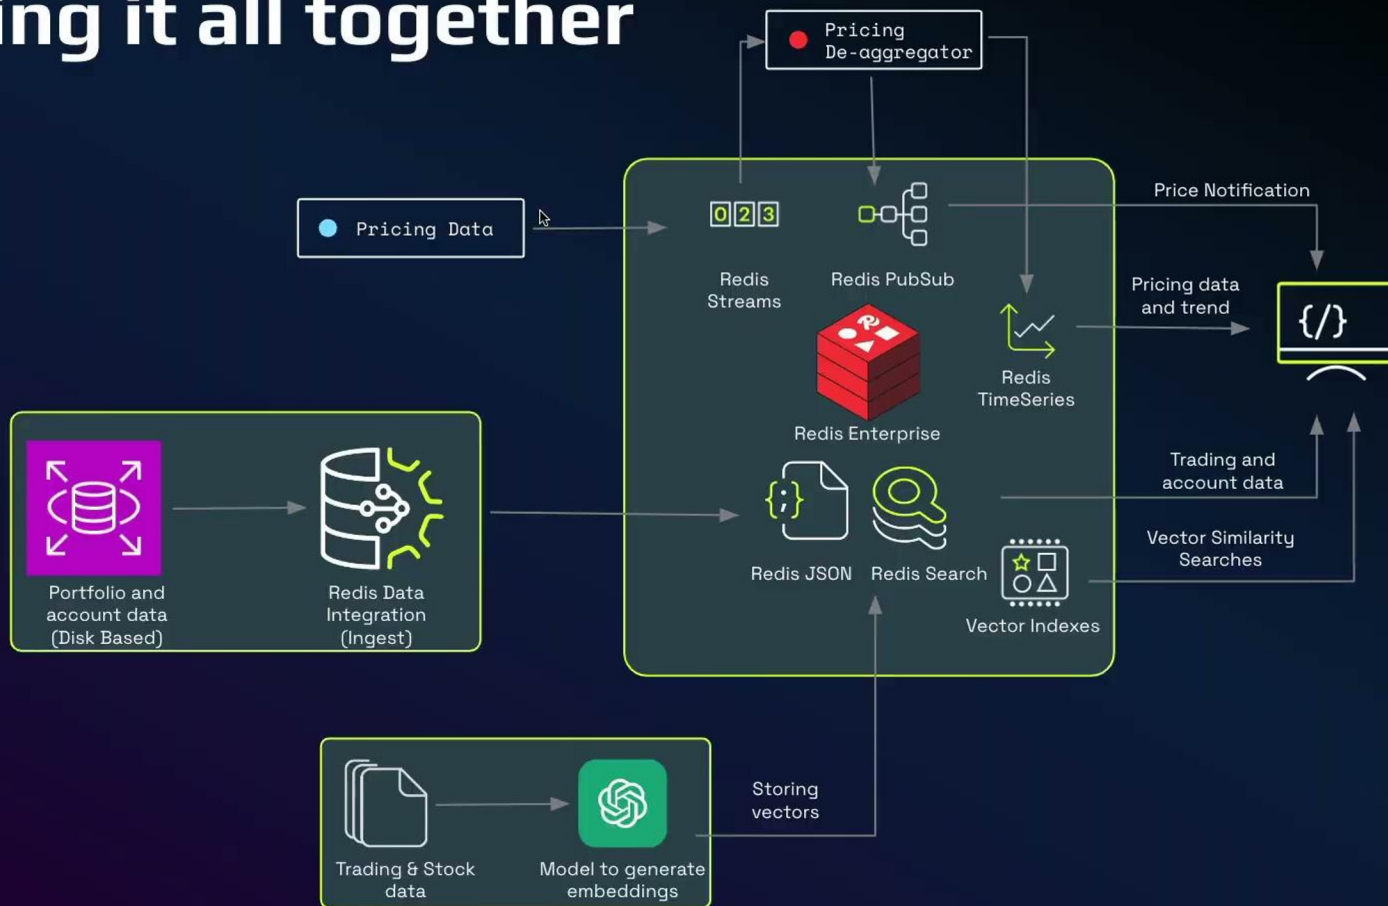
# Showcase

## A sample securities trading and account portfolio application

---



# Putting it all together



# Thank you!



Please complete the session  
survey in the mobile app